Pooling Pyramid Network for Object Detection

Pengchong Jin   Vivek Rathod   Xiangxin Zhu

Google AI Perception

{pengchong, rathodv, xiangxin}@google.com

Abstract

We'd like share a simple tweak of the SSD family detectors, which is effective in reducing the model size while maintaining the same quality. We share the predictors across all the scales, and replace the convolution between scales with max pooling. This has two advantages over the SSD: (1) avoids score miscalibration across scales; (2) the shared predictor sees the training data over all scale. Since we reduce the number of predictor to one, and trim all the convolutions between them, the model size is significantly reduced. We empirically found that it didn't hurt the model quality compared with its SSD counterpart.

1. Introduction

The SSD family detectors [5, 3] have been popular as they run fast, are simple to implement and easily portable to different types of hardware.

Most of the SSD detectors have several feature maps representing different scales, each of which uses its own predictor to produce the boxes and class scores. In practice, especially when the data distribution is skewed over scales, this design could run into problem. Imagine a dataset with tons of large objects and very few small ones. Those predictors from the small scale feature maps will be wasted as they rarely see any positives. This data imbalance could also result in score miscalibration across scale even for the same class. Another issue under this design is that each predictor only sees the objects at its own scale. This partitions will divide the already small dataset into even smaller pieces. If we believe the object appearance is scale invariant, it will be a more efficient way of using the data if all the predictors see all the data.

We propose simple changes to the SSD: use the same predictor in all scales. In order for the predictor to work in the same feature space, we replace the convolutions between feature maps with max pooling.

2. Pooling Pyramid Network (PPN)

The proposed model, *Pooling Pyramid Network (PPN)*, is a single-stage convolutional object detector, very similar to SSD with simple changes. The prediction head is designed to be light-weighted, fast to run, while maintains the comparable detection accuracy with its SSD counterpart. The network architecture is illustrated in Figure 1. There are two major changes to the original SSD [5]: (1) the box predictor is shared across feature maps with different scales; (2) the convolutions between feature maps are replaced with the max pooling operations. In the following sections, we will discuss the rationales behind them and effects of these changes.

2.1. Shared Box Predictor

The SSD uses independent box predictors for feature maps at different scales. The one potential problem is miscalibration of the prediction scores across different scales.

Since each box predictor is trained independently using only a portion of the groundtruth boxes that it is assigned to, different box predictors could see very different amount of positive and negative examples during the training. This implicit data imbalance could cause the problem that scores from different predictors fall in vastly different ranges, which makes them incomparable and difficult to use in the subsequent score-based postprocessing such as non maximum suppression. We design PPN with a shared box predictor across feature maps of different scales. As a result, the box predictor sees all the training data during the training when there are imbalance groundtruth boxes with different scales. This reduces the effect of miscalibration and unstable prediction scores.

One could argue that having separate box predictor for each scale increases the total capacity, and allows each predictor to focus at its specific scale. However, we think hat this may not be necessary as objects are mostly scale invariant. In practice, Faster-RCNN [6] works well with a single shared predictor.
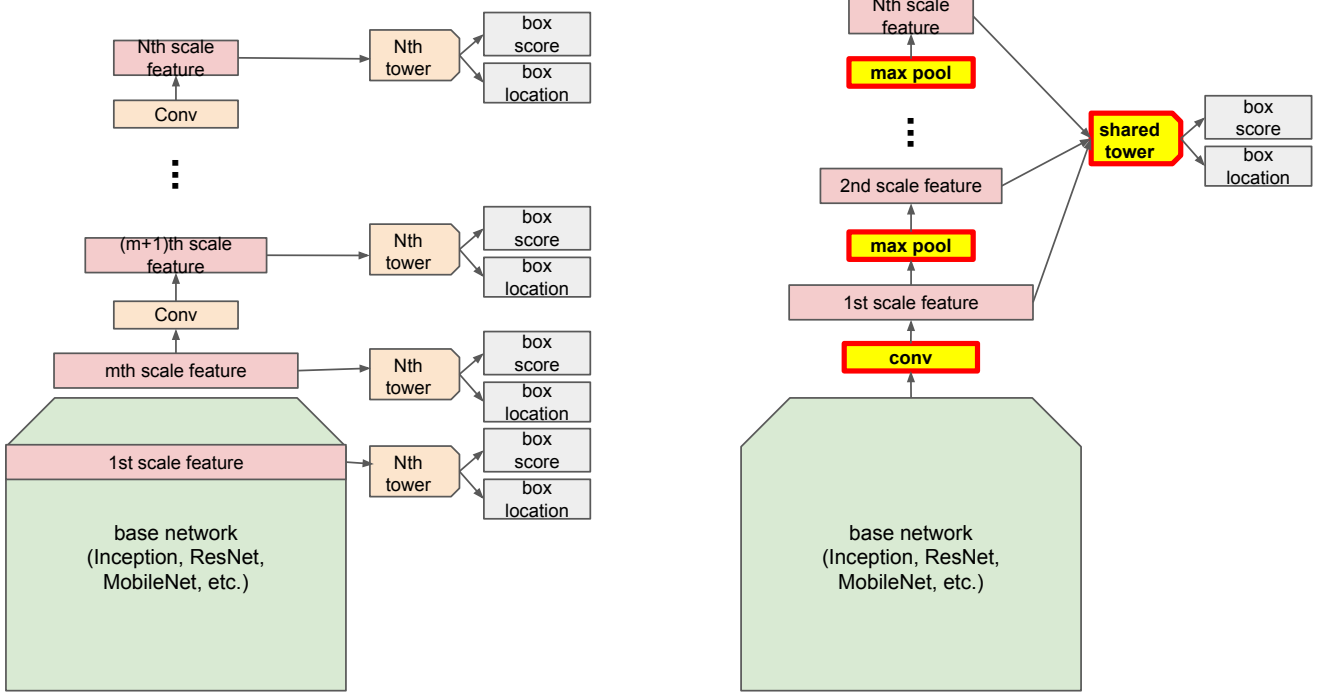
Figure 1. The architecture comparison between the Pooling Pyramid Network (PPN) and SSD. Left: SSD, Right: PPN. Note that the changes in PPN are highlighted: (1) using max pool to build the feature pyramid, (2) using the shared convolutional predictor for box classification and regression.

2.2. Max Pooling Pyramid

Our goal is to build a multi-scale feature pyramid structure, from which we can make the predictions using the shared box predictor. We achieve this by shrinking down a base feature map from the backbone network several times using a series of max pooling operations. This is different from SSD where feature maps are built by extracting layers from backbone network and shrinking them using additional convolutions, and FPN where feature maps are built by a top-down pathway with skip connections. We choose max pooling mainly for two reasons. First, using the pooling operations ensures feature maps with different scales live in the same embedding space, which makes training the shared box predictor more effective. In addition, since max pooling does not require any additions and multiplicatons, it is very fast to compute during the inference, therefore, making it suitable for many latency sensitive applications.

One may use the seemingly more intuitive avearge pooling to build the feature pyramid. It should be noted, however, that nonlinear operations after pooling need to be inserted, because otherwise it is a linear operations so that scores from the lower level feature maps would be higher than those from the pooled ones. We haven't experimented with the average pooling yet.

2.3. Overall Architecture

The final network architecture of our Pooling Pyramid Network (PPN) detector is illustrated in Figure 1. Followed by the backbone network, an optional 1×1 convolution is used to transform the features from the backbone network to a space with desired dimensions. We then apply a series of stride-2 max pooling operations to shrink the feature map down to 1×1. A shared box predictor is applied to feature maps of different scales in order to produce classification scores and location offsets of box predictions. We add one additional shared convolution in the box predictor after pooling operations to prepare the feature to be used for predictions.

3. Experiments

We run the experiments on COCO [4] detection dataset and compare the performance of PPN with SSD. We use MobileNet v1 [1] as the backbone network and set the input resolution to be 300×300. Both models use the standard implementation of MobileNet-v1 SSD in Google Object Detection API [2]. For PPN, we extract the layer *Conv2d_11_pointwise* as the base feature map, from which we build 6 pooled feature maps that are of sizes 19×19, 10×10, 5×5, 3×3, 2×2, and 1×1. A shared 1×1

Model	mAP	inference FLOPs	number of parameters	GPU inference time
MobileNet SSD	20.0	2.48B	6.83M	27ms
MobileNet PPN	19.7	2.35B	2.18M	26ms

Table 1. COCO detection: MobileNet SSD vs MobileNet PPN

depth 512 convolution is applied before the box classifier and location regressor. We use the similar anchor design as SSD, the smooth l_1 loss for box regression, and the focal loss with $\alpha = 0.25$ and $\gamma = 2$ for box classification. Our implementation is based on Google Object Detection API and it is publicly available under Tensorflow's Github repository.

Both SSD and PPN models are initialized using the MobileNet-v1 checkpoint that is pre-trained on ImageNet, and both of them are trained on `trainval35k` and tested on `minival`. We leverage Google Cloud TPU for fast training. We perform the model benchmark using an Nvidia GeForce GTX TITAN X card. Table 1 shows the comparison between SSD and PPN. PPN achieves the similar mAP (19.7 vs 20.0), comparable FLOPs and inference time, but 3x smaller in model size.

References

[1] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. *CoRR*, abs/1704.04861, 2017.

[2] J. Huang, V. Rathod, C. Sun, M. Zhu, A. Korattikara, A. Fathi, I. Fischer, Z. Wojna, Y. Song, S. Guadarrama, and K. Murphy. Speed/accuracy trade-offs for modern convolutional object detectors. In *CVPR*, 2017.

[3] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár. Focal loss for dense object detection. In *ICCV*, 2017.

[4] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perorna, D. Damanan, P. Dollár, and C. L. Zitnick. Microsoft COCO: Common objects in context. 2014.

[5] W. Liu, D. Anguelov, D. Erhan, C. Szegedy, S. Reed, C.-Y. Fu, and A. C. Berg. SSD: Single shot multibox detector. In *ECCV*, 2016.

[6] S. Ren, K. He, R. Girshick, and J. Sun. Faster R-CNN: Towards real-time object detection with region proposal networks. In *Neural Information Processing Systems (NIPS)*, 2015.