

Pooling Pyramid Network for Object Detection

Pengchong Jin

Vivek Rathod

Xiangxin Zhu

Google AI Perception

{pengchong, rathodv, xiangxin}@google.com

Abstract

We'd like to share a simple tweak of the Single Shot Multibox Detector (SSD) family of detectors, which is effective in reducing model size while maintaining the same quality. We share box predictors across all scales, and replace convolution between scales with max pooling. This has two advantages over vanilla SSD: (1) it avoids score miscalibration across scales; (2) the shared predictor sees the training data over all scales. Since we reduce the number of predictors to one, and trim all convolutions between them, model size is significantly smaller. We empirically show that these changes does not hurt model quality compared to vanilla SSD.

1. Introduction

SSD detectors [6, 4] have been popular as they run fast, are simple to implement and easily portable to different types of hardware.

Most SSD detectors have several feature maps representing different scales, each of which uses its own predictor to produce boxes and class scores. In practice, especially when the data distribution is skewed over scales, this design is problematic. Imagine a dataset with many large objects and very few small ones. The predictors from small scale feature maps will be wasted as they rarely see any positives. This data imbalance could also result in score miscalibration across scales even for the same class. Another issue with this design is that each predictor only sees the objects at its own scale. This partition will divide the already small dataset into even smaller sets. If we believe that object appearance is scale invariant, it will be more efficient if all the predictors see all of the data.

We propose simple changes to vanilla SSD: use the same predictor for all scales. In order for the predictor to work in the same feature space, we replace convolutions between feature maps with max pooling.

2. Pooling Pyramid Network (PPN)

The proposed model, *Pooling Pyramid Network (PPN)*, is a single-stage convolutional object detector, very similar to vanilla SSD with simple changes. The prediction head is designed to be light-weight, fast to run, while maintaining comparable detection accuracy with vanilla SSD. The network architecture is illustrated in Figure 1. There are two major changes to vanilla SSD [6]: (1) the box predictor is shared across feature maps with different scales; (2) the convolutions between feature maps are replaced with max pooling operations. In the following sections, we will discuss the rationale behind these changes and their effects.

2.1. Shared Box Predictor

Vanilla SSD uses independent box predictors for feature maps at different scales. One problem is miscalibration of the prediction scores across different scales.

Since each box predictor is trained independently using only a portion of the groundtruth boxes that it is assigned to, different box predictors could see very different number of positive and negative examples during training. This implicit data imbalance causes the problem that scores from different predictors fall in vastly different ranges, which makes them incomparable and difficult to use in subsequent score-based postprocessing steps such as non maximum suppression. We design PPN with a shared box predictor across feature maps of different scales. As a result, the box predictor sees all of the training data even when there is an imbalance in groundtruth box scales. This reduces the effect of miscalibration and unstable prediction scores.

One could argue that having a separate box predictor for each scale increases the total capacity, and allows each predictor to focus on its specific scale. However, we think that this may not be necessary as objects are mostly scale invariant.

2.2. Max Pooling Pyramid

Our goal is to build a multi-scale feature pyramid structure, from which we can make predictions using the shared box predictor. We achieve this by shrinking a base feature

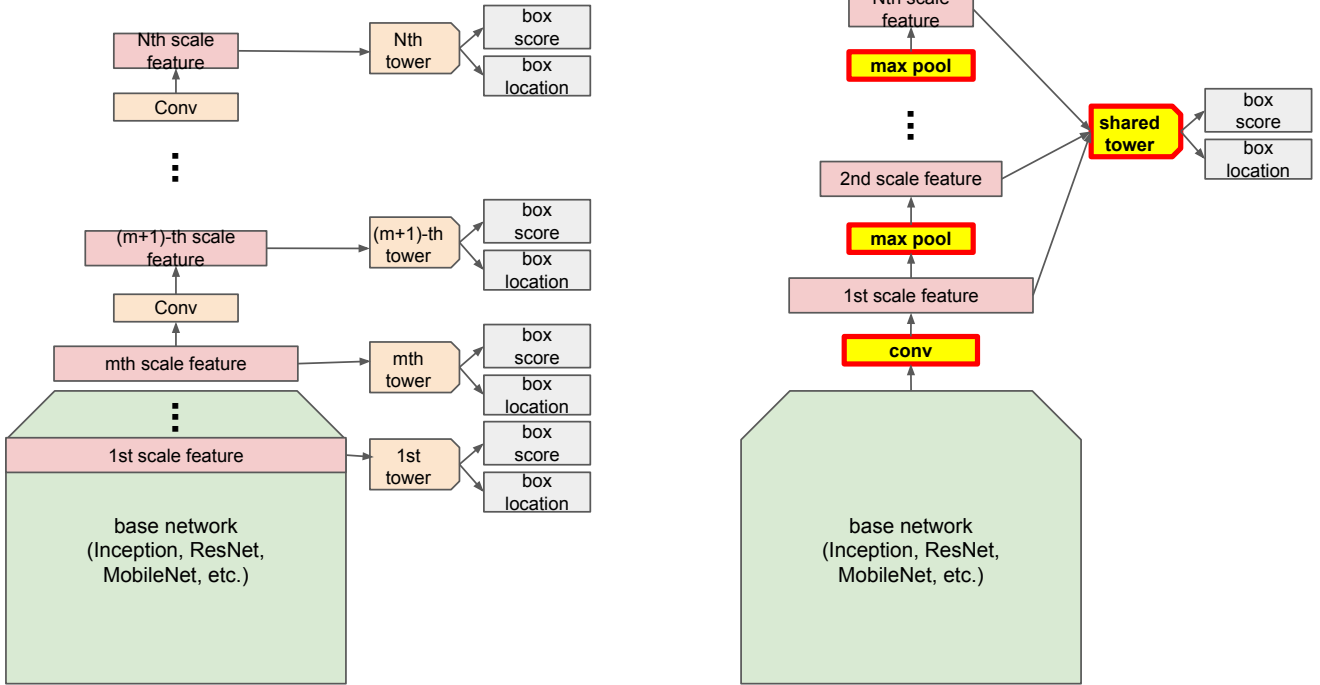


Figure 1. Architecture comparison between the Pooling Pyramid Network (PPN) and vanilla SSD. Left: vanilla SSD, Right: PPN. Note that the changes in PPN are highlighted: (1) using max pool to build the feature pyramid, (2) using shared convolutional predictors for box classification and regression.

map from the backbone network several times using a series of max pooling operations. This is different from vanilla SSD where feature maps are built by extracting layers from a backbone network and shrinking them using additional convolutions, and FPN where feature maps are built by a top-down pathway with skip connections. We choose max pooling mainly for two reasons. First, using the pooling operations ensures feature maps with different scales live in the same embedding space, which makes training the shared box predictor more effective. In addition, since max pooling does not require any additions and multiplications, it is very fast to compute during inference, making it suitable for many latency sensitive applications.

2.3. Overall Architecture

The final network architecture of our Pooling Pyramid Network (PPN) detector is illustrated in Figure 1. Followed by the backbone network, an optional 1×1 convolution is used to transform the features from the backbone network to a space with desired dimensions. We then apply a series of stride-2 max pooling operations to shrink the feature map down to 1×1 . A shared box predictor is applied to feature maps of different scales in order to produce classification scores and location offsets of box predictions. We add one additional shared convolution in the box predictor

after pooling operations to prepare the feature to be used for predictions.

3. Experiments

We run experiments on the COCO [5] detection dataset and compare the performance of PPN with vanilla SSD. We use MobileNet v1 [1] as the backbone network and set the input resolution to be 300×300 . Both models use the standard implementation of MobileNet-v1 SSD in the Tensorflow Object Detection API [2]. For PPN, we extract the layer *Conv2d_11_pointwise* as the base feature map, from which we build 6 pooled feature maps that are of sizes 19×19 , 10×10 , 5×5 , 3×3 , 2×2 , and 1×1 . A shared 1×1 depth 512 convolution is applied before the box classifier and location regressor. We use the same anchor design as SSD, smooth l_1 loss for box regression, and focal loss with $\alpha = 0.25$ and $\gamma = 2$ for box classification [4]. Our implementation is based on the Tensorflow Object Detection API and is publicly available under Tensorflow’s Github repository.

Both SSD and PPN models are initialized using a MobileNet-v1 checkpoint that is pre-trained on ImageNet, and both of them are trained and tested on the splits described in [2]. We leverage TPUs [3] for fast training. We perform the model benchmark using an Nvidia GeForce

Model	mAP	inference FLOPs	number of parameters	GPU inference time
MobileNet SSD	20.8	2.48B	6.83M	27ms
MobileNet PPN	20.3	2.35B	2.18M	26ms

Table 1. COCO detection: MobileNet SSD vs MobileNet PPN

GTX TITAN X card. Table 1 shows the comparison between SSD and PPN. PPN achieves similar mAP (20.3 vs 20.8), comparable FLOPs and inference time, but is 3x smaller in model size.

References

- [1] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. *CoRR*, abs/1704.04861, 2017.
- [2] J. Huang, V. Rathod, C. Sun, M. Zhu, A. Korattikara, A. Fathi, I. Fischer, Z. Wojna, Y. Song, S. Guadarrama, and K. Murphy. Speed/accuracy trade-offs for modern convolutional object detectors. In *CVPR*, 2017.
- [3] N. P. Jouppi, C. Young, N. Patil, D. Patterson, G. Agrawal, R. Bajwa, S. Bates, S. Bhatia, N. Boden, A. Borchers, R. Boyle, P.-I. Cantin, C. Chao, C. Clark, J. Coriell, M. Daley, M. Dau, J. Dean, B. Gelb, T. V. Ghaemmamghami, R. Gottipati, W. Gulland, R. Hagmann, C. R. Ho, D. Hogberg, J. Hu, R. Hundt, D. Hurt, J. Ibarz, A. Jaffey, A. Jaworski, A. Kaplan, H. Khaitan, D. Killebrew, A. Koch, N. Kumar, S. Lacy, J. Laudon, J. Law, D. Le, C. Leary, Z. Liu, K. Lucke, A. Lundin, G. MacKean, A. Maggiore, M. Mahony, K. Miller, R. Nagarajan, R. Narayanaswami, R. Ni, K. Nix, T. Norrie, M. Omernick, N. Penukonda, A. Phelps, J. Ross, M. Ross, A. Salek, E. Samadiani, C. Severn, G. Sizikov, M. Snellham, J. Souter, D. Steinberg, A. Swing, M. Tan, G. Thorson, B. Tian, H. Toma, E. Tuttle, V. Vasudevan, R. Walter, W. Wang, E. Wilcox, and D. H. Yoon. In-datacenter performance analysis of a tensor processing unit. In *Proceedings of the 44th Annual International Symposium on Computer Architecture, ISCA '17*, pages 1–12, New York, NY, USA, 2017. ACM.
- [4] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár. Focal loss for dense object detection. In *ICCV*, 2017.
- [5] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Damanan, P. Dollár, and C. L. Zitnick. Microsoft COCO: Common objects in context. 2014.
- [6] W. Liu, D. Anguelov, D. Erhan, C. Szegedy, S. Reed, C.-Y. Fu, and A. C. Berg. SSD: Single shot multibox detector. In *ECCV*, 2016.